

# Project Wine

Yongfan Zhao

2024-02-15

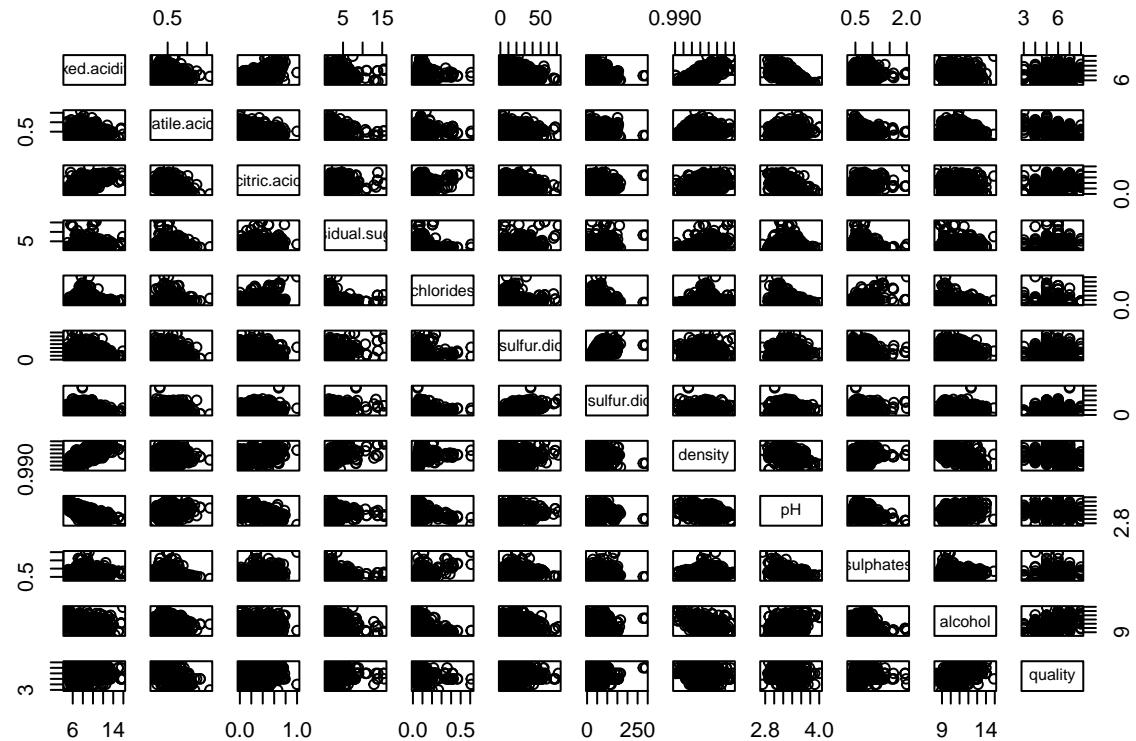
```
rm(list = ls())
Red_Wine = read.csv("winequality-red.csv")
View(Red_Wine)
Red_Wine = na.omit(Red_Wine)
# use summary function to understand the Data
# range, median, mean
summary(Red_Wine)

##   fixed.acidity  volatile.acidity  citric.acid  residual.sugar
##   Min.    : 4.60  Min.    :0.1200  Min.    :0.000  Min.    : 0.900
##   1st Qu.: 7.10  1st Qu.:0.3900  1st Qu.:0.090  1st Qu.: 1.900
##   Median  : 7.90  Median  :0.5200  Median  :0.260  Median  : 2.200
##   Mean    : 8.32  Mean    :0.5278  Mean    :0.271  Mean    : 2.539
##   3rd Qu.: 9.20  3rd Qu.:0.6400  3rd Qu.:0.420  3rd Qu.: 2.600
##   Max.    :15.90  Max.    :1.5800  Max.    :1.000  Max.    :15.500
##   chlorides      free.sulfur.dioxide total.sulfur.dioxide  density
##   Min.    :0.01200  Min.    : 1.00  Min.    : 6.00  Min.    :0.9901
##   1st Qu.:0.07000  1st Qu.: 7.00  1st Qu.:22.00  1st Qu.:0.9956
##   Median  :0.07900  Median  :14.00  Median  :38.00  Median  :0.9968
##   Mean    :0.08747  Mean    :15.87  Mean    :46.47  Mean    :0.9967
##   3rd Qu.:0.09000  3rd Qu.:21.00  3rd Qu.:62.00  3rd Qu.:0.9978
##   Max.    :0.61100  Max.    :72.00  Max.    :289.00  Max.    :1.0037
##   pH          sulphates      alcohol      quality
##   Min.    :2.740  Min.    :0.3300  Min.    : 8.40  Min.    :3.000
##   1st Qu.:3.210  1st Qu.:0.5500  1st Qu.: 9.50  1st Qu.:5.000
##   Median  :3.310  Median  :0.6200  Median  :10.20  Median  :6.000
##   Mean    :3.311  Mean    :0.6581  Mean    :10.42  Mean    :5.636
##   3rd Qu.:3.400  3rd Qu.:0.7300  3rd Qu.:11.10  3rd Qu.:6.000
##   Max.    :4.010  Max.    :2.0000  Max.    :14.90  Max.    :8.000

# sd for the each predictor
sd = sapply(Red_Wine, sd)
sd

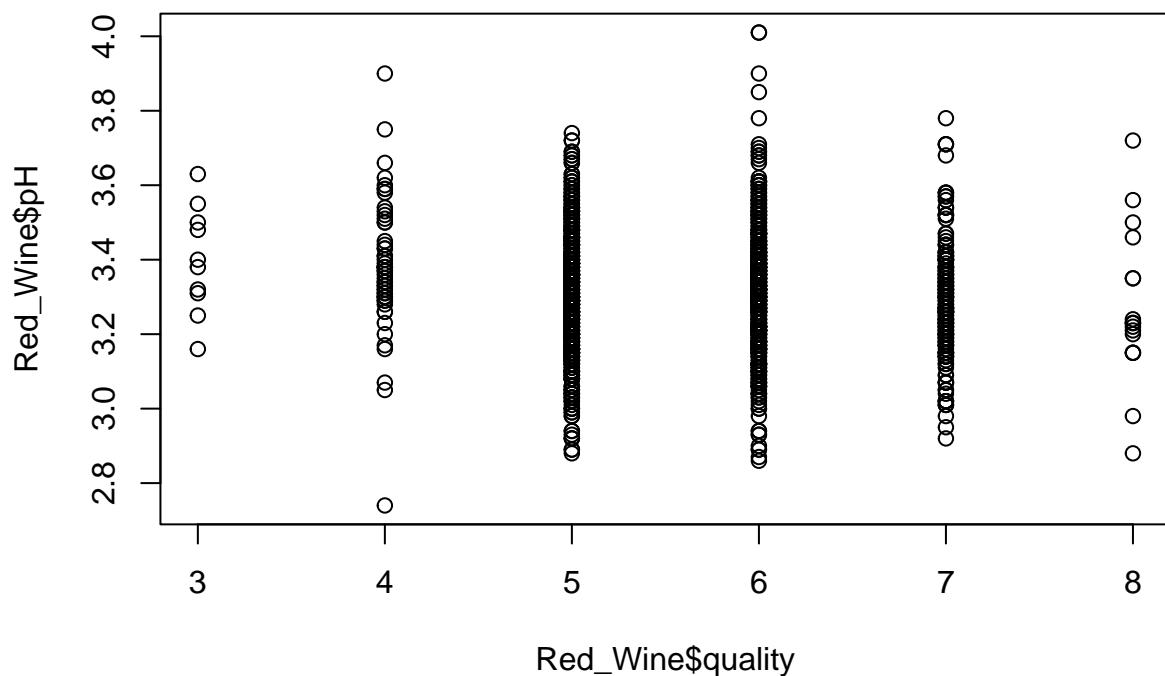
##       fixed.acidity  volatile.acidity  citric.acid
##       1.741096318  0.179059704  0.194801137
##       residual.sugar  chlorides  free.sulfur.dioxide
##       1.409928060  0.047065302  10.460156970
##       total.sulfur.dioxide  density      pH
##       32.895324478  0.001887334  0.154386465
##       sulphates      alcohol      quality
##       0.169506980  1.065667582  0.807569440
```

```
# use pairs to see the relation between target and predictor
pairs(Red_Wine)
```

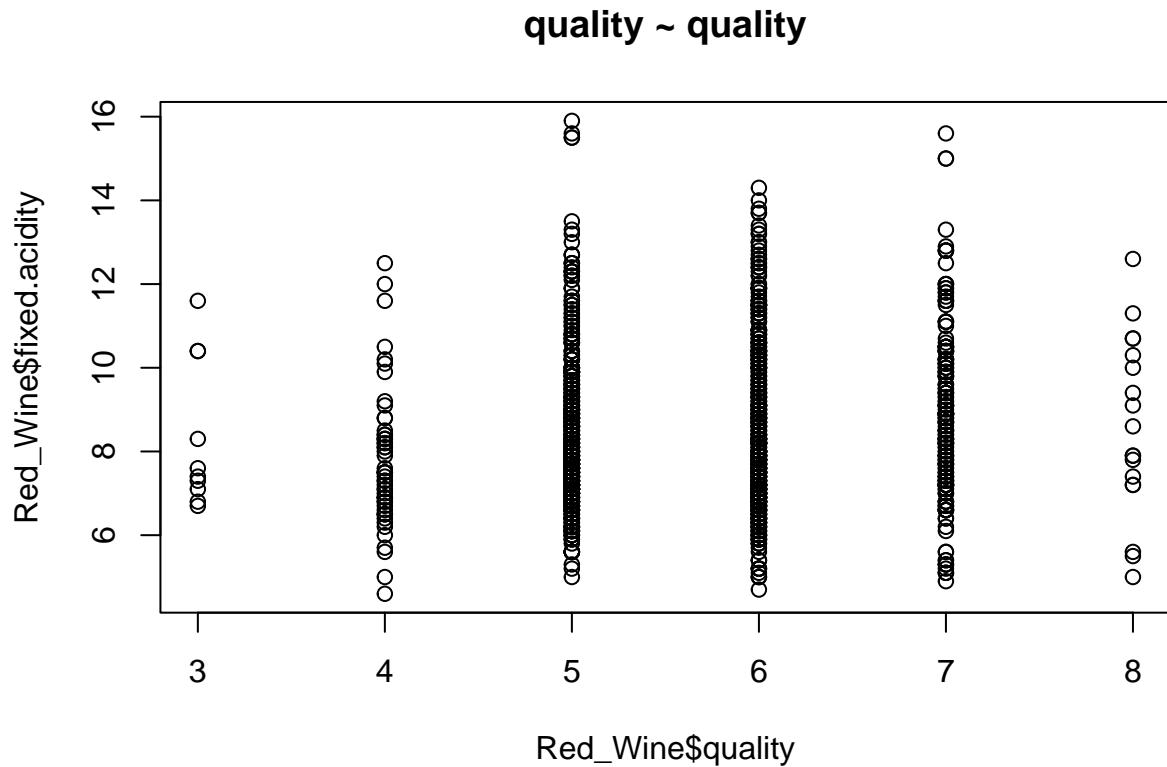


```
plot(Red_Wine$quality, Red_Wine$pH, main = "quality ~ pH")
```

**quality ~ pH**



```
plot(Red_Wine$quality, Red_Wine$fixed.acidity, main = "quality ~ quality")
```



```
# correlation of each data
```

```
round(cor(Red_Wine), digits = 3)
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar
## fixed.acidity	1.000	-0.256	0.672	0.115
## volatile.acidity	-0.256	1.000	-0.552	0.002
## citric.acid	0.672	-0.552	1.000	0.144
## residual.sugar	0.115	0.002	0.144	1.000
## chlorides	0.094	0.061	0.204	0.056
## free.sulfur.dioxide	-0.154	-0.011	-0.061	0.187
## total.sulfur.dioxide	-0.113	0.076	0.036	0.203
## density	0.668	0.022	0.365	0.355
## pH	-0.683	0.235	-0.542	-0.086
## sulphates	0.183	-0.261	0.313	0.006
## alcohol	-0.062	-0.202	0.110	0.042
## quality	0.124	-0.391	0.226	0.014
	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density
## fixed.acidity	0.094	-0.154	-0.113	0.668
## volatile.acidity	0.061	-0.011	0.076	0.022
## citric.acid	0.204	-0.061	0.036	0.365
## residual.sugar	0.056	0.187	0.203	0.355
## chlorides	1.000	0.006	0.047	0.201
## free.sulfur.dioxide	0.006	1.000	0.668	-0.022
## total.sulfur.dioxide	0.047	0.668	1.000	0.071
## density	0.201	-0.022	0.071	1.000
## pH	-0.265	0.070	-0.066	-0.342
## sulphates	0.371	0.052	0.043	0.149
## alcohol	-0.221	-0.069	-0.206	-0.496

```

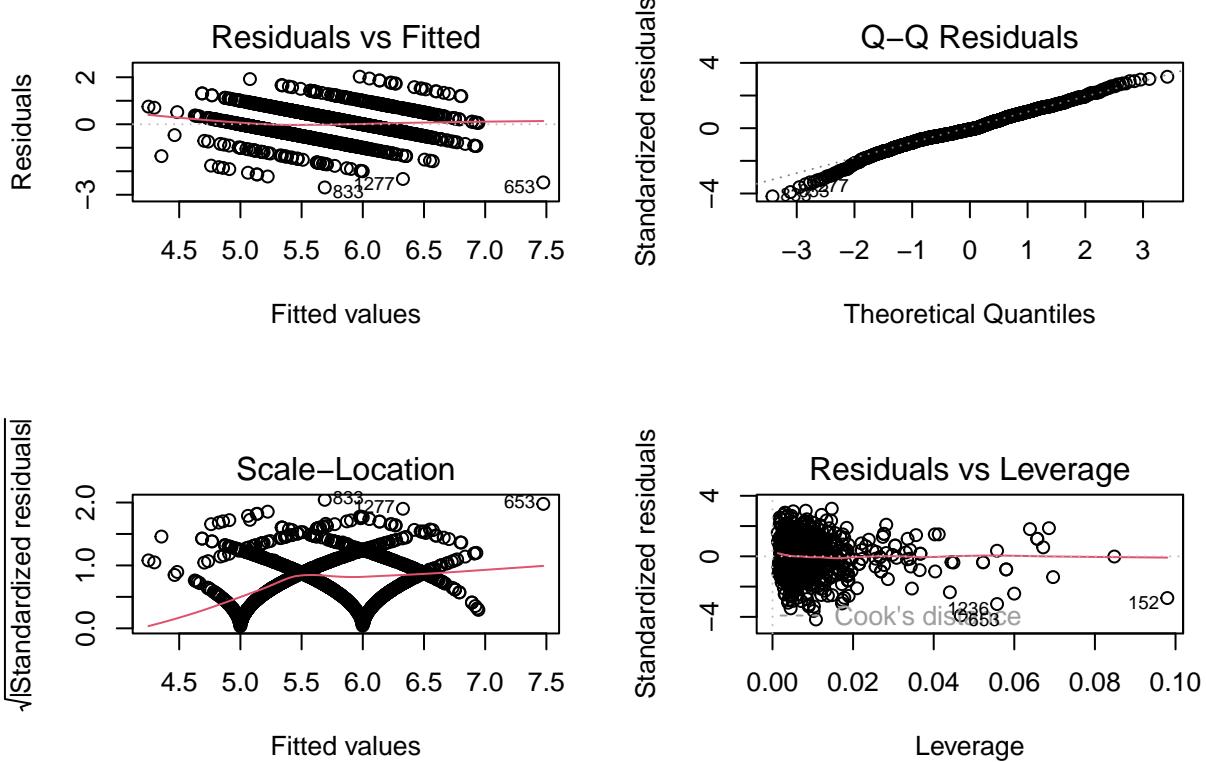
## quality          -0.129          -0.051          -0.185  -0.175
##               pH sulphates alcohol quality
## fixed.acidity   -0.683      0.183  -0.062  0.124
## volatile.acidity 0.235     -0.261  -0.202  -0.391
## citric.acid    -0.542      0.313  0.110  0.226
## residual.sugar -0.086      0.006  0.042  0.014
## chlorides       -0.265      0.371  -0.221  -0.129
## free.sulfur.dioxide 0.070      0.052  -0.069  -0.051
## total.sulfur.dioxide -0.066      0.043  -0.206  -0.185
## density         -0.342      0.149  -0.496  -0.175
## pH              1.000     -0.197  0.206  -0.058
## sulphates       -0.197      1.000  0.094  0.251
## alcohol          0.206      0.094  1.000  0.476
## quality         -0.058      0.251  0.476  1.000

lm.fit = lm(Red_Wine$quality ~ ., data = Red_Wine)
summary(lm.fit)

##
## Call:
## lm(formula = Red_Wine$quality ~ ., data = Red_Wine)
##
## Residuals:
##      Min      1Q      Median      3Q      Max 
## -2.68911 -0.36652 -0.04699  0.45202  2.02498 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 2.197e+01 2.119e+01  1.036  0.3002    
## fixed.acidity 2.499e-02 2.595e-02  0.963  0.3357    
## volatile.acidity -1.084e+00 1.211e-01 -8.948 < 2e-16 *** 
## citric.acid -1.826e-01 1.472e-01 -1.240  0.2150    
## residual.sugar 1.633e-02 1.500e-02  1.089  0.2765    
## chlorides    -1.874e+00 4.193e-01 -4.470 8.37e-06 *** 
## free.sulfur.dioxide 4.361e-03 2.171e-03  2.009  0.0447 *  
## total.sulfur.dioxide -3.265e-03 7.287e-04 -4.480 8.00e-06 *** 
## density      -1.788e+01 2.163e+01 -0.827  0.4086    
## pH            -4.137e-01 1.916e-01 -2.159  0.0310 *  
## sulphates    9.163e-01 1.143e-01  8.014 2.13e-15 *** 
## alcohol       2.762e-01 2.648e-02 10.429 < 2e-16 *** 
## ---    
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 0.648 on 1587 degrees of freedom
## Multiple R-squared:  0.3606, Adjusted R-squared:  0.3561 
## F-statistic: 81.35 on 11 and 1587 DF, p-value: < 2.2e-16

par(mfrow = c(2,2))
plot(lm.fit)

```



```
#####
set.seed(66)
dim(Red_Wine)

## [1] 1599   12
attach(Red_Wine)
train = sample(nrow(Red_Wine), nrow(Red_Wine)/2, replace = FALSE)
lm.fit1 = lm(quality ~ alcohol + volatile.acidity + sulphates+total.sulfur.dioxide + chlorides, data = Red_Wine)
summary(lm.fit1)

##
## Call:
## lm(formula = quality ~ alcohol + volatile.acidity + sulphates +
##     total.sulfur.dioxide + chlorides, data = Red_Wine[train,
## ])
##
## Residuals:
##      Min        1Q    Median        3Q       Max
## -2.25853 -0.40285 -0.04007  0.42492  2.00348
##
## Coefficients:
## (Intercept)            3.3312044  0.2909791  11.448 < 2e-16 ***
## alcohol                  0.2556504  0.0238567  10.716 < 2e-16 ***
## volatile.acidity       -1.3134004  0.1378472  -9.528 < 2e-16 ***
## sulphates                 0.9728511  0.1617876   6.013 2.78e-09 ***
## total.sulfur.dioxide   -0.0025558  0.0007723  -3.309 0.000978 ***
## chlorides                 -2.1740042  0.5345435  -4.067 5.24e-05 ***
## ---
```

```

## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6612 on 793 degrees of freedom
## Multiple R-squared:  0.356, Adjusted R-squared:  0.3519
## F-statistic: 87.67 on 5 and 793 DF, p-value: < 2.2e-16
Wine_test = Red_Wine[-train,"quality"]
lm.pred = predict(lm.fit1,Red_Wine[-train,])
lm_mse = mean((lm.pred -Wine_test)^2)

# linear regression with interaction
lm.fit2 = lm(quality ~ I(alcohol^2) + I(volatile.acidity^2) +log(sulphates)+log(total.sulfur.dioxide) +
summary(lm.fit2)

##
## Call:
## lm(formula = quality ~ I(alcohol^2) + I(volatile.acidity^2) +
##      log(sulphates) + log(total.sulfur.dioxide) + log(chlorides),
##      data = Red_Wine[train, ])
##
## Residuals:
##       Min     1Q     Median     3Q     Max 
## -2.33255 -0.41172 -0.04979  0.44163  2.06029
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)               4.702351   0.247212 19.022 < 2e-16 ***
## I(alcohol^2)              0.011465   0.001122 10.215 < 2e-16 ***
## I(volatile.acidity^2)     -0.998813   0.109468 -9.124 < 2e-16 ***
## log(sulphates)            0.818717   0.117289  6.980 6.23e-12 ***
## log(total.sulfur.dioxide) -0.107535   0.034411 -3.125 0.001843 ** 
## log(chlorides)             -0.293377   0.078544 -3.735 0.000201 *** 
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6589 on 793 degrees of freedom
## Multiple R-squared:  0.3605, Adjusted R-squared:  0.3564
## F-statistic: 89.39 on 5 and 793 DF, p-value: < 2.2e-16
lm.pred2 = predict(lm.fit2,Red_Wine[-train,])
lm1_mse = mean((lm.pred2 -Wine_test)^2)

```

Ridge Regression

```

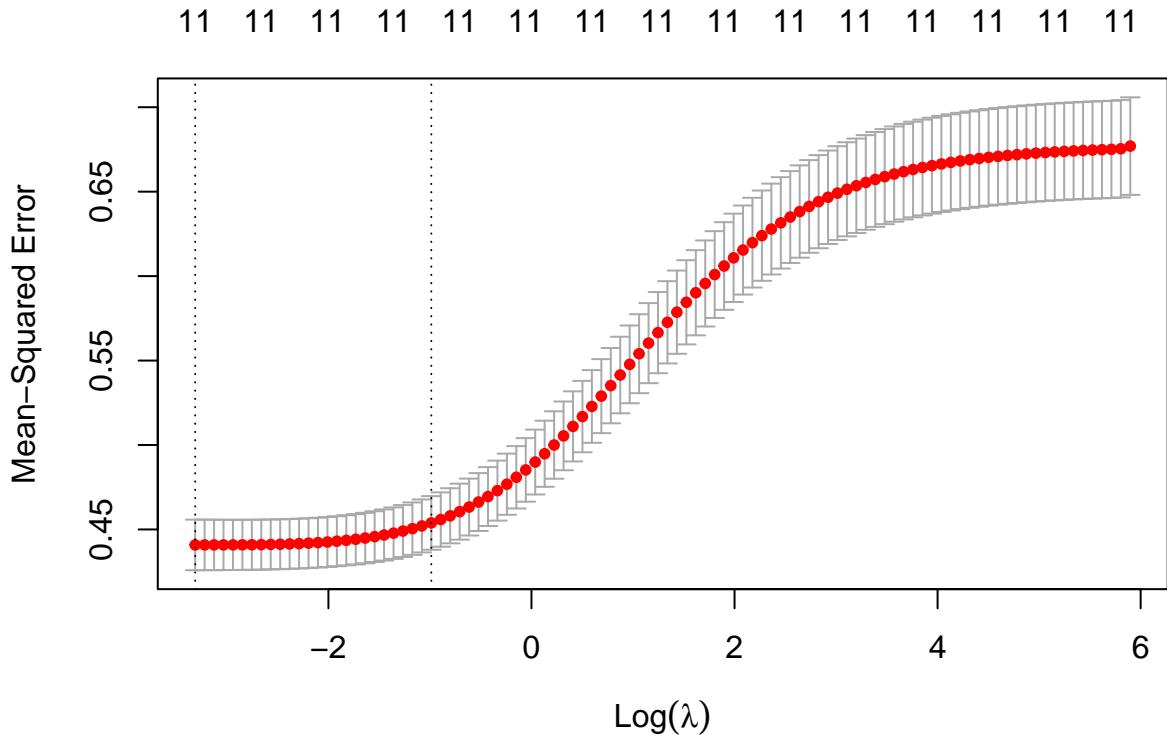
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-8

x = model.matrix(quality~.,Red_Wine)[,-1]
y = Red_Wine$quality
grid = 10 ^ seq(10,-2,length = 100)
ridge.mod = glmnet(x[train,],y[train],alpha = 0,lambda = grid,thresh = 1e-12)
# find the best lambda
cv.out = cv.glmnet(x[train,],y[train],alpha = 0)
plot(cv.out)

```



```

bestlam = cv.out$lambda.min
bestlam

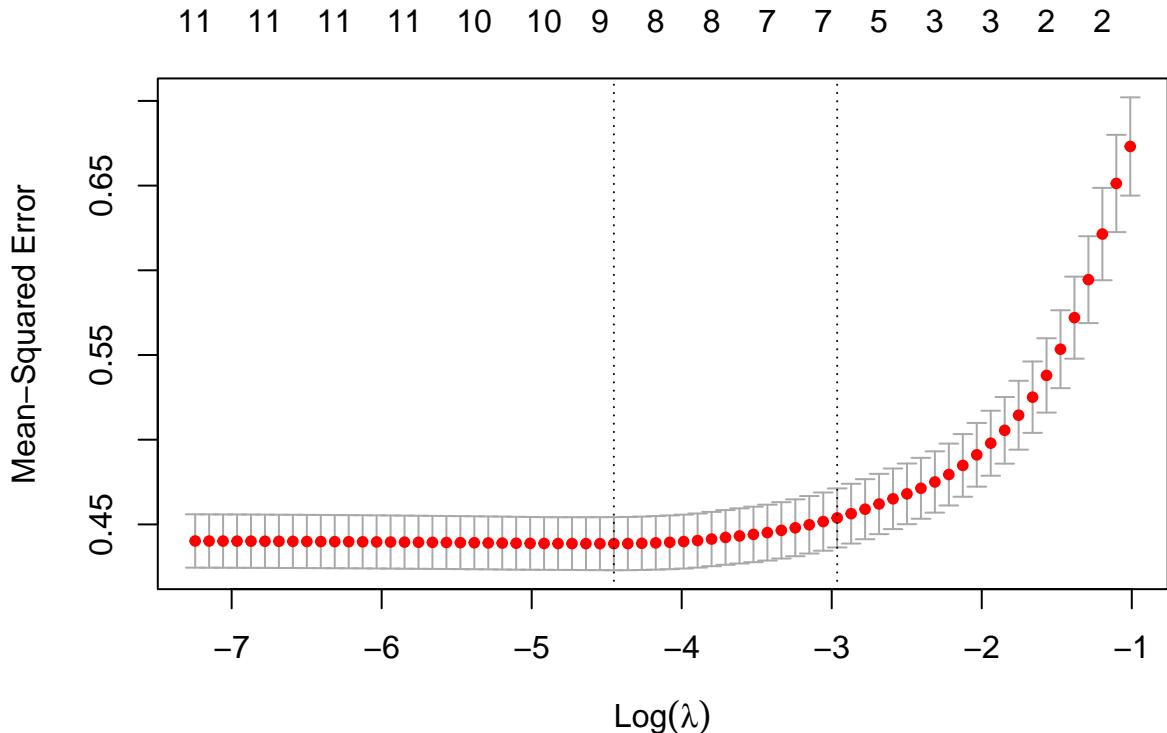
## [1] 0.03642799

ridge.pred = predict(ridge.mod,s = bestlam,x = x[train,],y = y[train],newx = x[-train,],exact = T)
ridge_mse = mean((ridge.pred-Wine_test)^2)

lasso Regression

lasso.mod = glmnet(x[train,],y[train],alpha = 1,lambda = grid,thresh = 1e-12)
# find the best lambda
cv.out.lasso = cv.glmnet(x[train,],y[train],alpha = 1)
plot(cv.out.lasso)

```



```

bestlam_lasso = cv.out.lasso$lambda.min
bestlam_lasso

## [1] 0.01165429

lasso.pred = predict(lasso.mod,s = bestlam,x = x[train,],y = y[train],newx = x[-train,],exact = T)
lasso_mse = mean((lasso.pred-Wine_test)^2)

lasso.coef = predict(lasso.mod,type = "coefficients",s = bestlam)[-1,]
# get non-zero coefficients
lasso.coef[lasso.coef!=0]

##      fixed.acidity      volatile.acidity      chlorides
##      0.004429025     -1.196743499    -1.193591994
## total.sulfur.dioxide          pH      sulphates
##      -0.001526670     -0.134384945      0.629998435
##      alcohol
##      0.251486602

```

for the non-zero coef, there are 7 coef, compare those three regression.  $R^2$ :Ridge < LM.FIT1 < LASSO < LM.FIT

```

# TREES
library(tree)
tree.wine = tree(quality ~ .,data = Red_Wine[train,])
summary(tree.wine)

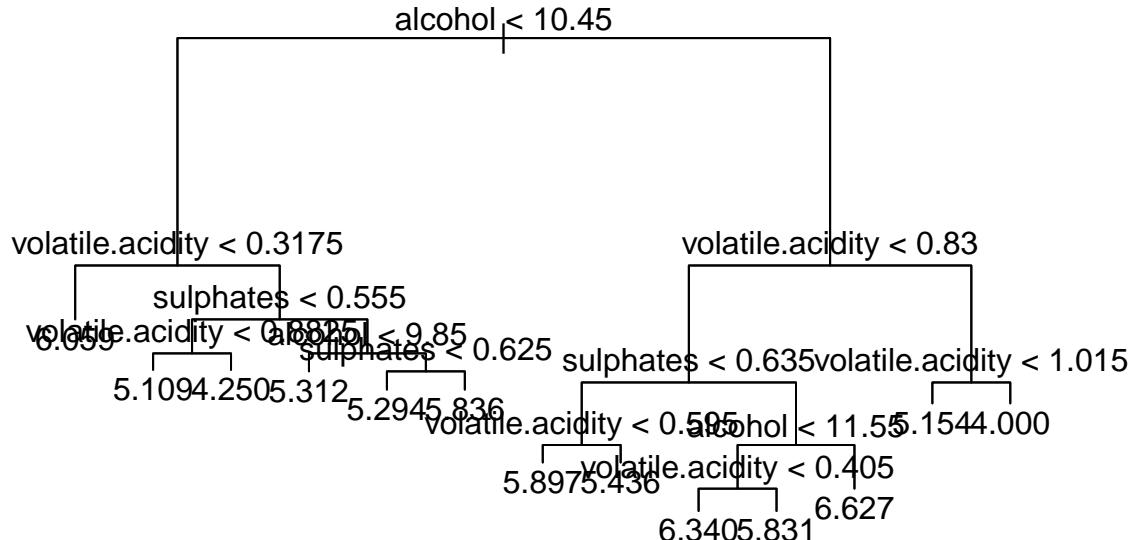
##
## Regression tree:
## tree(formula = quality ~ ., data = Red_Wine[train, ])
## Variables actually used in tree construction:
## [1] "alcohol"           "volatile.acidity" "sulphates"
## Number of terminal nodes: 13

```

```

## Residual mean deviance: 0.4044 = 317.8 / 786
## Distribution of residuals:
##   Min. 1st Qu. Median Mean 3rd Qu. Max.
## -2.3120 -0.3122 -0.1094 0.0000 0.3733 1.9410
# 13 nodes,
#plot the tree
plot(tree.wine)
text(tree.wine,pretty = 0)

```



```

tree.pred = predict(tree.wine,Red_Wine[-train,])
tree_mse = mean((tree.pred - Wine_test)^2)

```

prune tree

```

cv.out.tree = cv.tree(tree.wine,FUN = prune.tree)
cv.out.tree

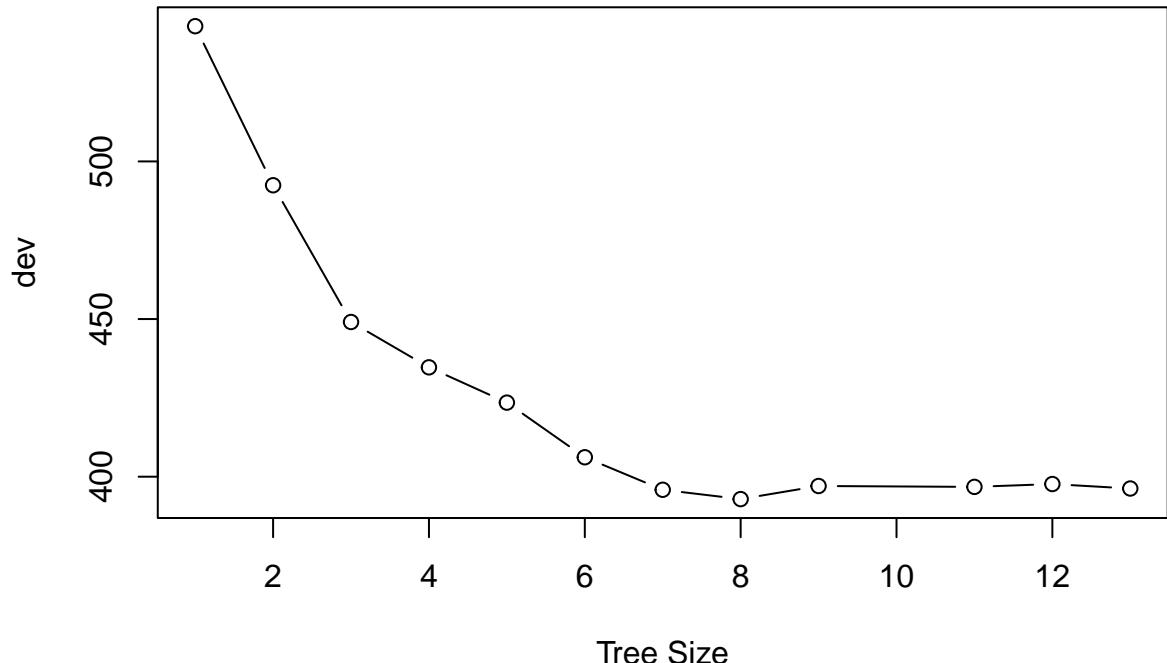
```

```

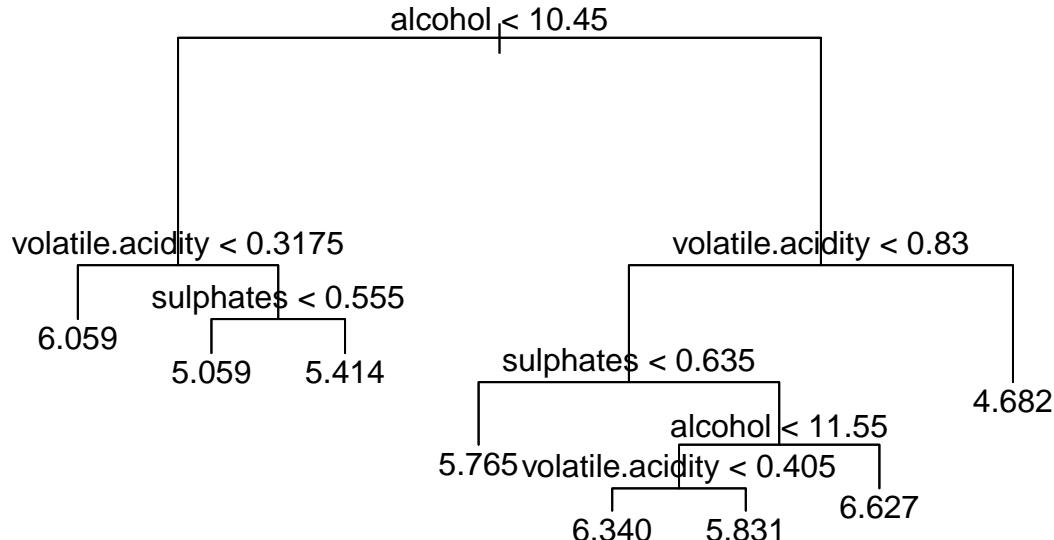
## $size
## [1] 13 12 11 9 8 7 6 5 4 3 2 1
##
## $dev
## [1] 396.2448 397.6669 396.7684 397.0445 392.8938 395.8512 406.1593 423.4958
## [9] 434.6871 449.0610 492.4407 542.8909
##
## $k
## [1]      -Inf  5.560662  5.911772  6.129975  7.080420  7.559492 11.499052
## [8] 14.759779 18.230322 21.167095 39.558796 76.913333
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune"           "tree.sequence"

```

```
plot(cv.out.tree$size, cv.out.tree$dev, type = "b", xlab = "Tree Size", ylab = "dev")
```



```
wine.prune = prune.tree(tree.wine, best = 8)
plot(wine.prune)
text(wine.prune, pretty = 0)
```



```
prune.pred = predict(wine.prune, Red_Wine[-train])
prune_mse = mean((prune.pred - Wine_test)^2)
```

```
## Warning in prune.pred - Wine_test: longer object length is not a multiple of
## shorter object length
```

```
# prune not work will in this case
```

Bagging

```

library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
bagging.fit = randomForest(quality~., data = Red_Wine, subset = train, mtry = 11, importance = TRUE)
bagging.fit

##
## Call:
##   randomForest(formula = quality ~ ., data = Red_Wine, mtry = 11,           importance = TRUE, subset = tra
##                 Type of random forest: regression
##                           Number of trees: 500
## No. of variables tried at each split: 11
##
##               Mean of squared residuals: 0.4006737
##                         % Var explained: 40.53

bagging.pred = predict(bagging.fit, Red_Wine[-train,])
bagging_mse = mean((bagging.pred-Wine_test)^2)

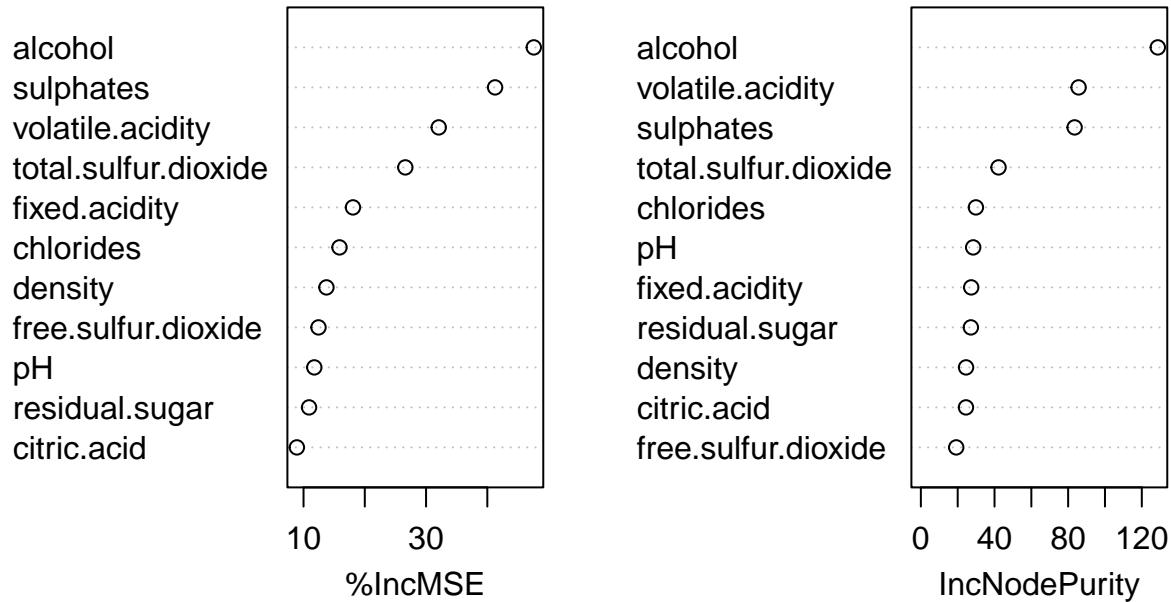
# random forest
min_mean = Inf
for (i in 1:11){
  rf.fit = randomForest(quality~., data = Red_Wine, subset = train, mtry = i, importance = TRUE)
  bagging.pred = predict(rf.fit, Red_Wine[-train,])
  rf_MSE = mean((bagging.pred-Wine_test)^2)
  if(rf_MSE < min_mean){
    min_mean = rf_MSE
    best_mtry = i
  }
}
cat("Minimum mean:", min_mean, "with mtry =", best_mtry, "\n")

## Minimum mean: 0.3342029 with mtry = 1
rf_mse = min_mean

varImpPlot(rf.fit)

```

## rf.fit



```
importance(rf.fit)
```

```
##          %IncMSE IncNodePurity
## fixed.acidity    18.071975    27.34916
## volatile.acidity 32.069219    85.67413
## citric.acid      8.912563    24.49292
## residual.sugar   10.902992   27.15757
## chlorides        15.872484   29.89000
## free.sulfur.dioxide 12.440197  19.21082
## total.sulfur.dioxide 26.613665  42.18373
## density          13.741017   24.53474
## pH                11.760270   28.40956
## sulphates        41.264372   83.50861
## alcohol           47.583639  128.68512
```

```
boosting
```

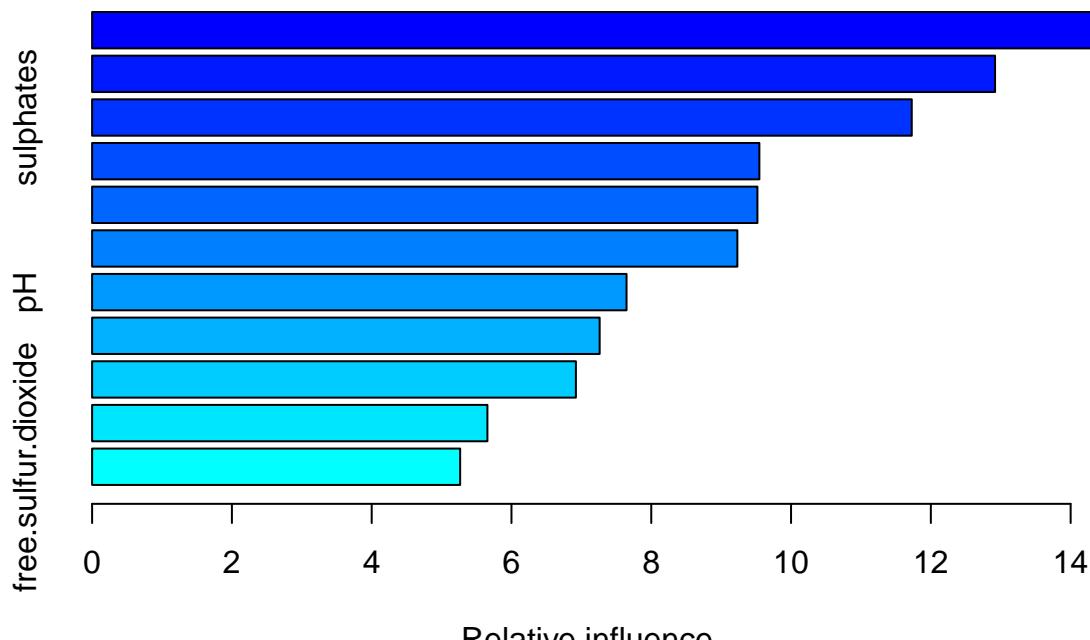
```
library(gbm)
```

```
## Loaded gbm 2.1.9
```

```
## This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com
```

```
boost.fit = gbm(quality~., data = Red_Wine[train,], distribution = "gaussian", n.trees = 5000, interaction.depth = 3, shrinkage = 0.01, n.minobsinnode = 10)
```

```
summary(boost.fit)
```



```

##                                var   rel.inf
## volatile.acidity      volatile.acidity 14.306817
## alcohol                  alcohol 12.918694
## sulphates                sulphates 11.726890
## total.sulfur.dioxide total.sulfur.dioxide  9.546487
## chlorides                 chlorides  9.518603
## density                   density  9.231772
## pH                         pH  7.646158
## citric.acid              citric.acid 7.261388
## fixed.acidity             fixed.acidity 6.921577
## residual.sugar            residual.sugar 5.655800
## free.sulfur.dioxide     free.sulfur.dioxide 5.265815

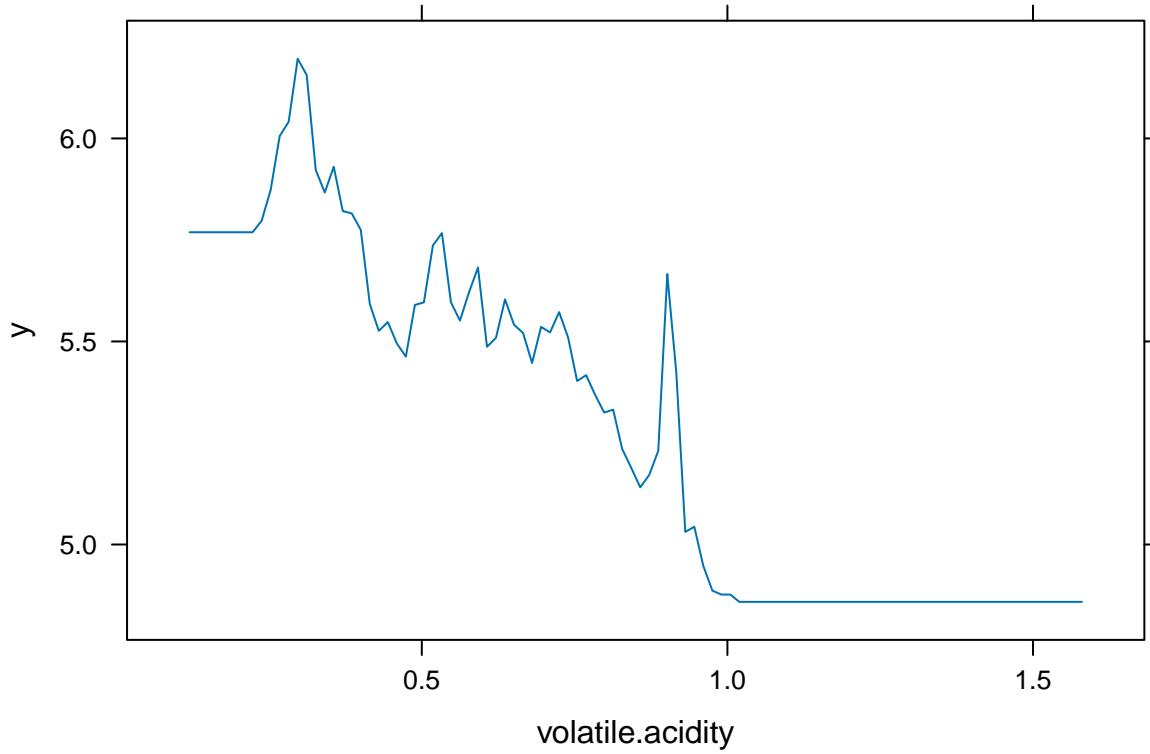
boost.pred = predict(boost.fit,newdata = Red_Wine[-train,],n.trees=5000)
mean((boost.pred-Wine_test)^2)

## [1] 0.4381224

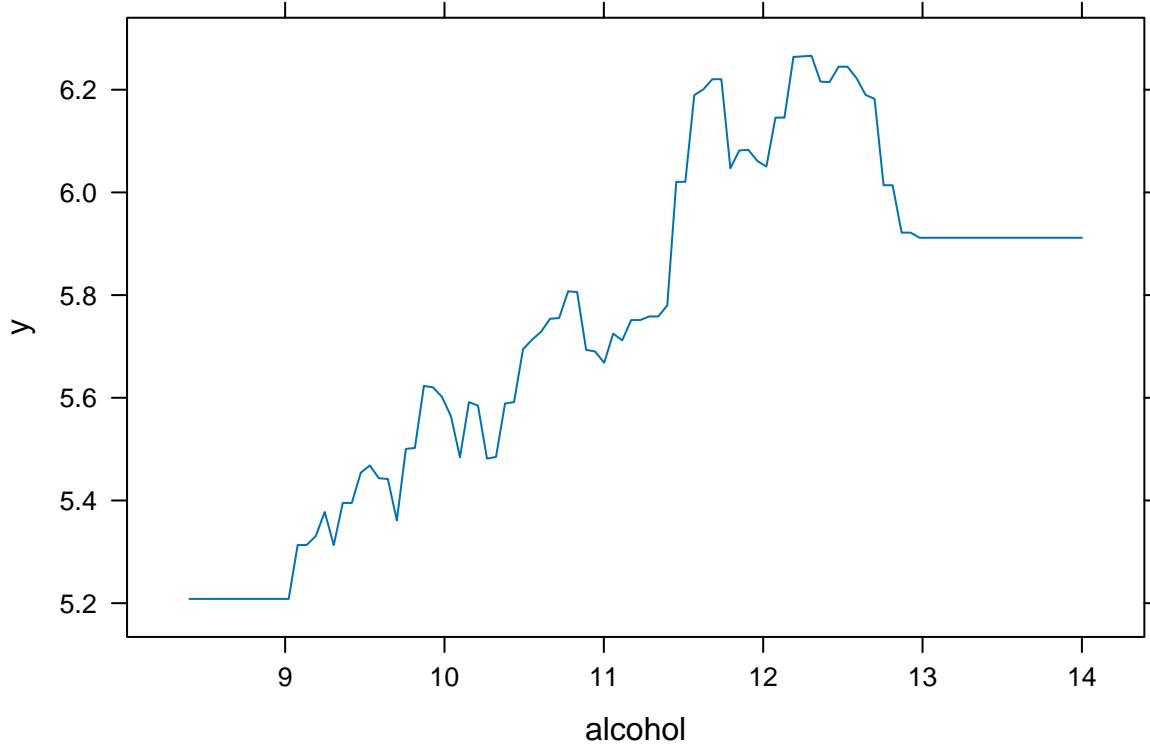
boost.pred = predict(boost.fit,newdata = Red_Wine[-train,],n.trees=500)
boost_mse = mean((boost.pred-Wine_test)^2)
# when we use less tree get a small MSE

plot(boost.fit,i = "volatile.acidity")

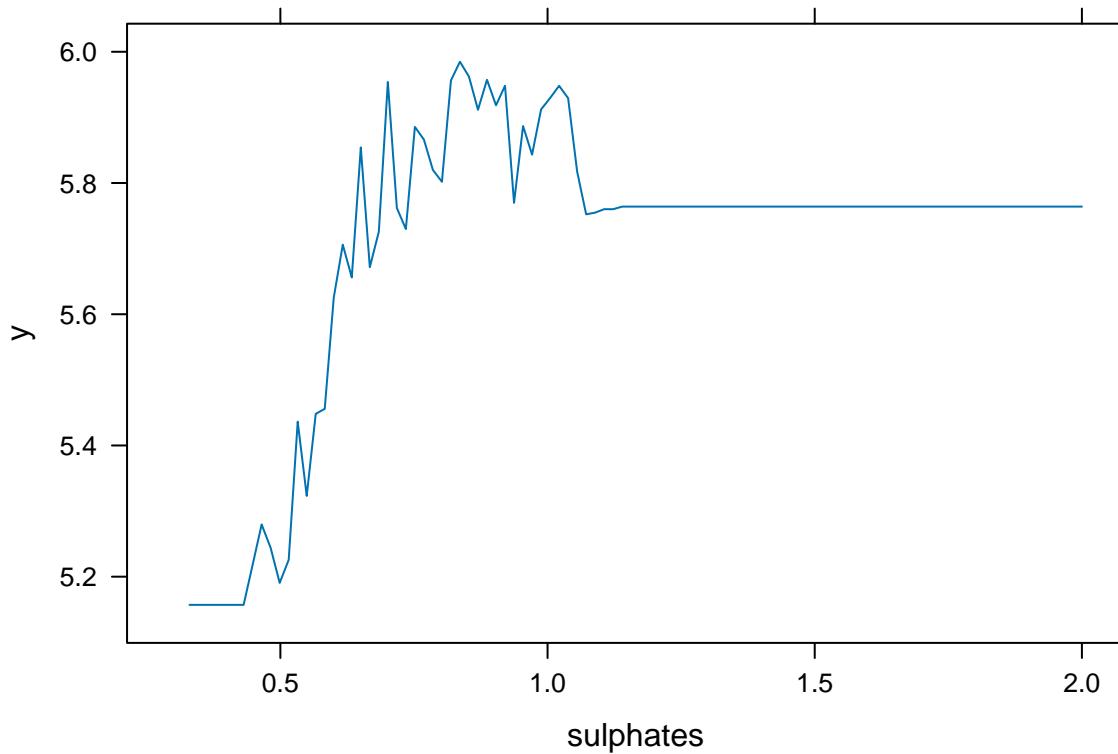
```



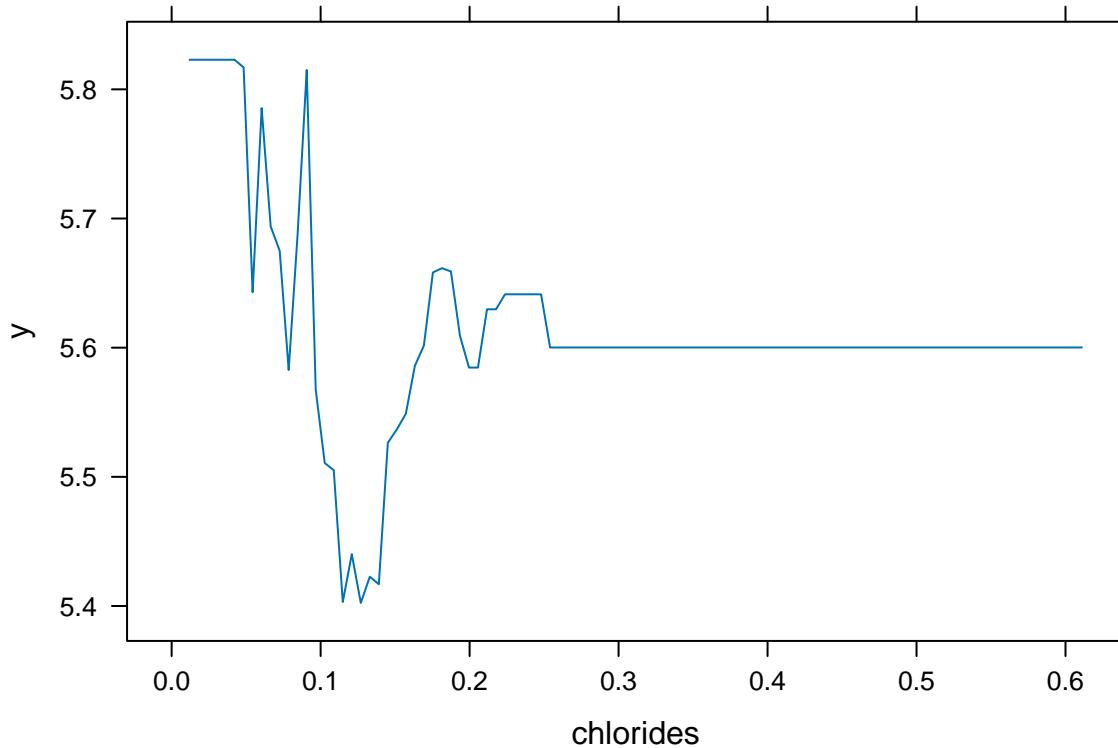
```
plot(boost.fit,i = "alcohol")
```



```
plot(boost.fit,i = "sulphates")
```



```
plot(boost.fit,i = "chlorides")
```



```
library(neuralnet)
library(caret)
```

```
## Loading required package: ggplot2
```

```

## 
## Attaching package: 'ggplot2'
## The following object is masked from 'package:randomForest':
## 
##     margin

## Loading required package: lattice
library(nnet)
df = read.csv("winequality-red.csv")
dim(df)

## [1] 1599   12
head(df)

##   fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1          7.4           0.70      0.00          1.9       0.076
## 2          7.8           0.88      0.00          2.6       0.098
## 3          7.8           0.76      0.04          2.3       0.092
## 4         11.2           0.28      0.56          1.9       0.075
## 5          7.4           0.70      0.00          1.9       0.076
## 6          7.4           0.66      0.00          1.8       0.075
##   free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates alcohol
## 1                 11            34 0.9978 3.51      0.56      9.4
## 2                 25            67 0.9968 3.20      0.68      9.8
## 3                 15            54 0.9970 3.26      0.65      9.8
## 4                 17            60 0.9980 3.16      0.58      9.8
## 5                 11            34 0.9978 3.51      0.56      9.4
## 6                 13            40 0.9978 3.51      0.56      9.4
##   quality
## 1      5
## 2      5
## 3      5
## 4      6
## 5      5
## 6      5

var = c("volatile.acidity","alcohol","sulphates","total.sulfur.dioxide","chlorides","density","pH","citric.acid")
df = df[,c('quality',var)]
str(df)

## 'data.frame': 1599 obs. of 9 variables:
## $ quality : int 5 5 5 6 5 5 5 7 7 5 ...
## $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
## $ alcohol : num 9.4 9.8 9.8 9.8 9.4 9.4 9.4 10 9.5 10.5 ...
## $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57 0.8 ...
## $ total.sulfur.dioxide: num 34 67 54 60 34 40 59 21 18 102 ...
## $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
## $ density : num 0.998 0.997 0.997 0.998 0.998 ...
## $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
## $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...

# make change in different dataset
#record the range & value of price
quality = df[, 'quality']
max_quality = range(df['quality'])[2]

```

```

min_quality = range(df['quality'])[1]
# scale the numerical
numerical = c("quality","volatile.acidity","alcohol","sulphates","total.sulfur.dioxide","chlorides","de"

norm.values = preProcess(df[,numerical],method = 'range')
df[,numerical] = predict(norm.values,df[,numerical])
head(df[,numerical])

##   quality volatile.acidity   alcohol sulphates total.sulfur.dioxide chlorides
## 1      0.4       0.3972603 0.1538462 0.1377246      0.09893993 0.1068447
## 2      0.4       0.5205479 0.2153846 0.2095808      0.21554770 0.1435726
## 3      0.4       0.4383562 0.2153846 0.1916168      0.16961131 0.1335559
## 4      0.6       0.1095890 0.2153846 0.1497006      0.19081272 0.1051753
## 5      0.4       0.3972603 0.1538462 0.1377246      0.09893993 0.1068447
## 6      0.4       0.3698630 0.1538462 0.1377246      0.12014134 0.1051753
##   density      pH citric.acid
## 1 0.5675477 0.6062992      0.00
## 2 0.4941263 0.3622047      0.00
## 3 0.5088106 0.4094488      0.04
## 4 0.5822320 0.3307087      0.56
## 5 0.5675477 0.6062992      0.00
## 6 0.5675477 0.6062992      0.00

#make change in different dataset
# convert categorical to dummies

f = as.formula(paste('quality~',
                      paste(names(df)[!names(df) %in% c("quality")],
                            collapse = '+')))

f

## quality ~ volatile.acidity + alcohol + sulphates + total.sulfur.dioxide +
##         chlorides + density + pH + citric.acid

set.seed(66)
train = sample(nrow(df),nrow(df)/2)
nn = neuralnet(f,data = df[train,],hidden = 2)
compute_mse = function(nn,df,train,quality){
  # make prediction with train
  pred.train = compute(nn,subset(df[train,],select = -c(quality)))
  pred.train.orig = pred.train$net.result *(max_quality - min_quality) + min_quality
  train.mse = mean((quality[train]-pred.train.orig)^2)
  # test
  pred.test = compute(nn,subset(df[-train,],select = -c(quality)))
  pred.test.orig = pred.test$net.result *(max_quality - min_quality) + min_quality
  test.mse = sqrt(mean((quality[-train]-pred.test.orig)^2))
  # return
  mse = as.data.frame(rbind(train.mse,test.mse))
  return(mse)
}
mse = compute_mse(nn,df,train,quality)
mse

##          V1
## train.mse 0.3862409
## test.mse  0.6264340

```

single layer with 5 nodes

```
nn1 = neuralnet(f,data = df[train,],hidden =4)
mse1 = compute_mse(nn1,df,train,quality)
```

{two layers, 5 nodes in each layer}

```
nn2 = neuralnet(f,data = df[train,],hidden =c(5,5))
mse2 = compute_mse(nn2,df,train,quality)
```

2layer,4 nodes in each layer

```
nn3 = neuralnet(f,data = df[train,],hidden =c(4,4))
mse3 = compute_mse(nn3,df,train,quality)
```

```
nn4 = neuralnet(f,data = df[train,],hidden =c(4,4,4))
mse4 = compute_mse(nn4,df,train,quality)
```

```
nn_mse_df = cbind(mse,mse1,mse3,mse4,mse2)
names(nn_mse_df) = c('1_layer_2nodes',"1_layer_4_nodes","2_layer_4_nodes","3_layer_4_nodes", "2_layer_5_nodes")
nn_mse_df
```

```
##          1_layer_2nodes 1_layer_4_nodes 2_layer_4_nodes 3_layer_4_nodes
## train.mse      0.3862409      0.3875833      0.3586865      0.2990268
## test.mse       0.6264340      0.6195124      0.6469280      0.6746118
##          2_layer_5_nodes
## train.mse      0.3428813
## test.mse       0.8324282
```

```
regression_mse_df = as.data.frame(cbind(lm_mse,lm1_mse,ridge_mse,lasso_mse))
names(regression_mse_df) = c("lm_mse","lm1_mse","ridge_mse","lasso_mse")
regression_mse_df
```

```
##      lm_mse  lm1_mse ridge_mse lass0_mse
## 1 0.4146733 0.412079 0.4113106 0.4117974
```

```
trees_mse_df = as.data.frame(cbind(tree_mse,prune_mse,bagging_mse,rf_mse,boost_mse))
names(trees_mse_df) = c("tree_mse","prune_mse","bagging_mse","rf_mse","boost_mse")
trees_mse_df
```

```
##      tree_mse prune_mse bagging_mse      rf_mse boost_mse
## 1 0.4482091 0.8966937 0.3342029 0.3342029 0.3785012
```